

# ICS Homework 5

Cao Shengcao

## 4.47

- A. It shouldn't be difficult to write the C code. Here we use a simple array of 10 long integers to test the function.

```
#include <stdio.h>
long a[10] = { 7, 4, 2, 6, 9, 0, 5, 1, 8, 3 }; // test array
void bubble_p(long *data, long count)        // bubblesort using pointers
{
    long *p, *q, *r, last;
    for (last = count - 1; last > 0; --last)
    {
        r = data + last;                    // p < r instead of i < last
        for (p = data; p < r; ++p)
        {
            q = p + 1;
            if (*q < *p)                    // *q is data[i + 1], and *p is data[i]
            {
                long t = *q;
                *q = *p;
                *p = t;
            }
        }
    }
}
int main()
{
    int i;
    for (i = 0; i < 10; ++i)
        printf("%ld ", a[i]);
    printf("\n");                          // 7 4 2 6 9 0 5 1 8 3
    bubble_p(a, 10);
    for (i = 0; i < 10; ++i)
        printf("%ld ", a[i]);
    printf("\n");                          // 0 1 2 3 4 5 6 7 8 9
    return 0;
}
```

- B. Imitating the Y86-64 code from figure 4-7 (P252) and the C code above, we can also finish the Y86-64 code for this program. Be careful about how we implement for-loops in assembly language. For a `long*` type pointer, `++p` is actually adding 8 to its value. (It really took me some time to debug this program...)

```

.pos 0

init:
    irmovq stack, %rsp
    call   main
    halt

    .align 8
a:      # test array
    .quad  7
    .quad  4
    .quad  2
    .quad  6
    .quad  9
    .quad  0
    .quad  5
    .quad  1
    .quad  8
    .quad  3

bubble_p:
    irmovq 1, %rax
    subq   %rax, %rsi      # %rsi: last = count - 1
    jle    bubble_p_end1

bubble_p_loop1:
    rrmovq %rsi, %rax
    addq   %rax, %rax
    addq   %rax, %rax
    addq   %rax, %rax      # %rax: last * 8
    rrmovq %rdi, %r8
    addq   %rax, %r8        # %r8: r = data + last
    rrmovq %rdi, %rdx      # %rdx: p = data
    rrmovq %rdx, %rax
    subq   %r8, %rax
    jge    bubble_p_end2

bubble_p_loop2:
    irmovq 8, %rcx
    addq   %rdx, %rcx      # %rcx: q = p + 1
    mrmovq (%rdx), %r9     # %r9: *p
    mrmovq (%rcx), %r10    # %r10: *q
    rrmovq %r10, %rax
    subq   %r9, %rax
    jge    no_swap

swap:
    rmmovq %r10, (%rdx)
    rmmovq %r9, (%rcx)    # swap

no_swap:
    irmovq 8, %rax
    addq   %rax, %rdx      # ++p

```

```

    rrmovq %rdx, %rax
    subq   %r8, %rax
    jl     bubble_p_loop2

bubble_p_end2:
    irmovq 1, %rax
    subq   %rax, %rsi    # --last
    jg     bubble_p_loop1

bubble_p_end1:
    ret

main:
    irmovq 10, %rsi
    irmovq a, %rdi
    call   bubble_p
    ret

    .pos 0x200

stack:

```

The Y86-64 simulator can be found at [CSAPP's student site](https://csapp.cs.cmu.edu/4e/y86/). Running the two programs, we can see the correct results:

```

friedrich@ubuntu: ~/Documents/pre
friedrich@ubuntu:~/Documents/pre$ ./sim/misc/yas 4.47.y
friedrich@ubuntu:~/Documents/pre$ ./sim/misc/yis 4.47.yo
Stopped in 718 steps at PC = 0x13. Status 'HLT', CC Z=1 S=0 O=0
Changes to registers:
%rax: 0x0000000000000000 0x0000000000000001
%rcx: 0x0000000000000000 0x0000000000000020
%rdx: 0x0000000000000000 0x0000000000000020
%rsp: 0x0000000000000000 0x0000000000000020
%rdi: 0x0000000000000000 0x0000000000000018
%r8: 0x0000000000000000 0x0000000000000020
%r10: 0x0000000000000000 0x0000000000000001

Changes to memory:
0x0018: 0x0000000000000007 0x0000000000000000
0x0020: 0x0000000000000004 0x0000000000000001
0x0030: 0x0000000000000006 0x0000000000000003
0x0038: 0x0000000000000009 0x0000000000000004
0x0040: 0x0000000000000000 0x0000000000000005
0x0048: 0x0000000000000005 0x0000000000000006
0x0050: 0x0000000000000001 0x0000000000000007
0x0060: 0x0000000000000003 0x0000000000000009
0x01f0: 0x0000000000000000 0x0000000000000125
0x01f8: 0x0000000000000000 0x0000000000000013
friedrich@ubuntu:~/Documents/pre$

friedrich@ubuntu:~/Documents/pre$ gcc 4.47.c -o 4.47
friedrich@ubuntu:~/Documents/pre$ ./4.47
7 4 2 6 9 0 5 1 8 3
0 1 2 3 4 5 6 7 8 9
friedrich@ubuntu:~/Documents/pre$

```

You may want to look into the details about how our handwritten Y86-64 assembly code differs from the assembly code generated by GCC. Check this out in [4.47.c](#) [4.47.s](#) [4.47.y](#) [4.47.yo](#).

## 4.56

To implement this so called "BTFNT", we need to modify three parts of our HCL code.

- Fetch stage: Predict the next PC as BTFNT describes. If the prediction were wrong, we must select the correct branch. Note there are two cases.

```

## What address should instruction be fetched at
word f_pc = [
    # Mispredicted branch. Fetch at incremented PC
    M_icode == IJXX && M_ifun != UNCOND && M_valE < M_valA && !M_Cnd : M_valA;
    M_icode == IJXX && M_ifun != UNCOND && M_valE >= M_valA && M_Cnd : M_valE;
    # Completion of RET instruction
    W_icode == IRET : W_valM;
    # Default: Use predicted value of PC
    1 : F_predPC;
];

# Predict next value of PC
word f_predPC = [
    # BBTFNT: This is where you'll change the branch prediction rule
    f_icode == IJXX && f_ifun != UNCOND && f_valC < f_valP : f_valC;
    f_icode == IJXX && f_ifun != UNCOND && f_valC >= f_valP : f_valP;
    f_icode in { IJXX, ICALL } : f_valC;
    1 : f_valP;
];

```

- Execute stage: Pass `valC` by calculating `valE = valB + valA = 0 + valC`. (`valP` is merged into `valA`, so we don't need to treat it specially.)

```

## Select input A to ALU
word aluA = [
    E_icode in { IRRMOVQ, IOPQ } : E_valA;
    E_icode in { IIRMOVQ, IRMMOVQ, IMRMOVQ, IJXX } : E_valC;
    E_icode in { ICALL, IPUSHQ } : -8;
    E_icode in { IRET, IPOPQ } : 8;
    # Other instructions don't need ALU
];

## Select input B to ALU
word aluB = [
    E_icode in { IRMMOVQ, IMRMOVQ, IOPQ, ICALL,
                IPUSHQ, IRET, IPOPQ } : E_valB;
    E_icode in { IRRMOVQ, IIRMOVQ, IJXX } : 0;
    # Other instructions don't need ALU
];

```

- Controlling: Insert two bubbles almost as the same as before. The conditions are changed.

```

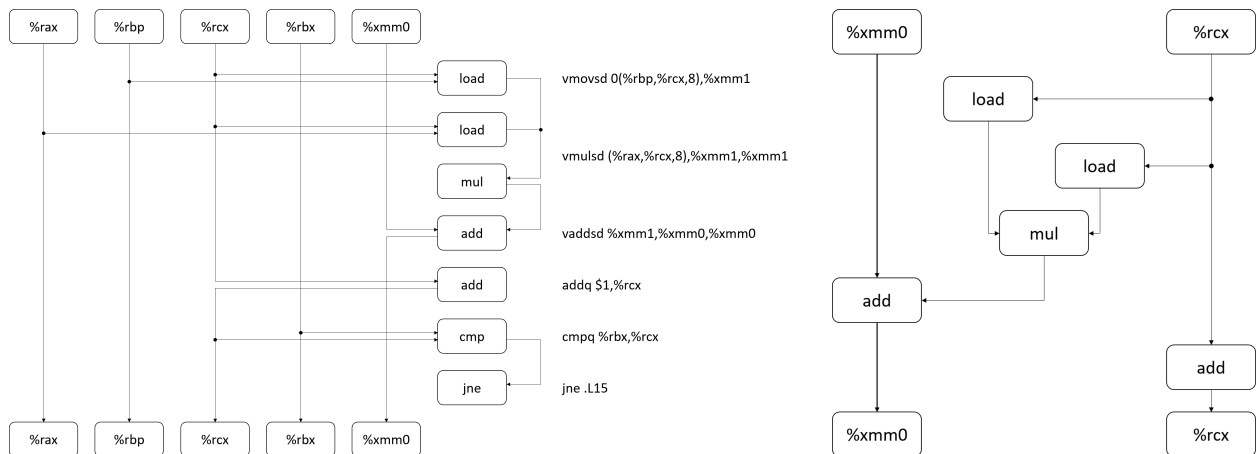
bool D_bubble =
    # Mispredicted branch
    (E_icode == IJXX && E_ifun != UNCOND && E_valC < E_valA && !e_Cnd) ||
    (E_icode == IJXX && E_ifun != UNCOND && E_valC >= E_valA && e_Cnd) ||
    # BBTFNT: This condition will change
    # Stalling at fetch while ret passes through pipeline
    # but not condition for a load/use hazard
    !(E_icode in { IMRMOVQ, IPOPOP } && E_dstM in { d_srcA, d_srcB }) &&
    IRET in { D_icode, E_icode, M_icode };

bool E_bubble =
    # Mispredicted branch
    (E_icode == IJXX && E_ifun != UNCOND && E_valC < E_valA && !e_Cnd) ||
    (E_icode == IJXX && E_ifun != UNCOND && E_valC >= E_valA && e_Cnd) ||
    # BBTFNT: This condition will change
    # Conditions for a load/use hazard
    E_icode in { IMRMOVQ, IPOPOP } &&
    E_dstM in { d_srcA, d_srcB};

```

## 5.13

- A. The diagrams are shown below.



- B. The lower bound is **3**, because the critical path only has one FP add instruction.
- C. The lower bound is **1**, as the critical path remains the same.
- D. The FP mul instructions can be pipelined. When the loop gets running, the results of FP mul instructions can be obtained in every cycle.

## 5.14

The code using  $6 \times 1$  loop unrolling:

```

void inner_loop_unrolling(vec_ptr u, vec_ptr v, data_t *dest)
{
    long i;
    long length = vec_length(u);
    long limit = length - 5;
    data_t *udata = get_vec_start(u);
    data_t *vdata = get_vec_start(v);
    data_t sum = (data_t)0;
    for (i = 0; i < limit; i += 6)
    {
        sum += udata[i] * vdata[i] + udata[i + 1] * vdata[i + 1]
              + udata[i + 2] * vdata[i + 2] + udata[i + 3] * vdata[i + 3]
              + udata[i + 4] * vdata[i + 4] + udata[i + 5] * vdata[i + 5];
    }
    for (; i < length; ++i)
    {
        sum += udata[i] * vdata[i];
    }
    *dest = sum;
}

```

- A. The performance is limited because two load instructions are required every time. Even if the two function units for load work all the time, the CPE can't be less than **1**.
- B. The processor has only one function unit for FP add, which is always in the critical path. Therefore, the CPE can't be less than **3**.